

# Assignment 2

## CS834 Introduction to Information Retrieval

### Fall 2017

Polykarpou Thomadakis

October 7, 2017

## Question 4.1

Plot rank-frequency curves (using a log-log graph) for words and bigrams in the Wikipedia collection available through the book website ([http://www.search- engines-book.com](http://www.search-engines-book.com)). Plot a curve for the combination of the two. What are the best values for the parameter c for each curve?

## Answer

To answer this question, I wrote a Python script that parses each documents in the Wikipedia collection and extracts all of its tokens. Then, count the frequency of each token, create bigrams from those tokens and count the frequency of the bigrams.

Finally, for both tokens and bigrams rank them by their frequency and compute their probability and the 'c' value. The output is written on a csv file for easier visualization into tables. For the purposes of this assignment I used the Wiki small collection of the book's website.

The script can be found in listing 1.

```
1 import os
2 import csv
3 from bs4 import BeautifulSoup
4 import nltk
5
6 # Set the path to search for html files
7 root_dir = 'en'
8
9 # Set the tokenizer expression
10 tokenizer = nltk.RegexpTokenizer(r'\w+')
11
12 # Find all files we need to parse
13 html_files = []
14 for root, dirs, files in os.walk(root_dir):
15     for file in files:
16         if file.endswith('.html'):
17             html_files.append(os.path.join(root, file))
18
19
20 # Get the tokens for all files
21 tokens_found = {}
22 num_words = 0
23 tokens_in_bigrams = {}
24 for file in html_files:
25     with open(file) as curr_file:
26         print file
27         soup = BeautifulSoup(curr_file.read(), 'html.parser')
28         tokens= tokenizer.tokenize(soup.get_text())
29         for token in tokens:
```

```

31     token = token.encode('utf-8')
32     if token not in tokens_found:
33         tokens_found[token] = 0
34     tokens_found[token] +=1
35     num_words +=1
36
37 # Get the bigrams from the tokens
38 for token in nltk.bigrams(tokens):
39     token = (token[0].encode('utf-8'),token[1].encode('utf-8'))
40     if token not in tokens_in_bigrams:
41         tokens_in_bigrams[token] = 0
42         tokens_in_bigrams[token] += 1
43
44 # Write token results to a file
45 f = open('results.csv','w+')
46 writer = csv.writer(f)
47 writer.writerow(['Rank','Frequency','Word','Probability','c'])
48 rank=0
49 for w in sorted(tokens_found, key=tokens_found.get, reverse=True):
50     rank+=1
51     writer.writerow( (rank,tokens_found[w],w,float(tokens_found[w])/num_words ,(
52         float(tokens_found[w])/num_words)*rank ))
53
54 f.close()
55
56 # Write bigram results to a file
57 f = open('results_bigrams.csv','w+')
58 writer = csv.writer(f)
59 writer.writerow(['Rank','Frequency','Word','Probability','c'])
60 rank=0
61 for w in sorted(tokens_in_bigrams, key=tokens_in_bigrams.get, reverse=True):
62     rank+=1
63     writer.writerow( (rank,tokens_in_bigrams[w],w,float(tokens_in_bigrams[w])/num_words ,(
64         float(tokens_in_bigrams[w])/num_words)*rank ))
65
66 f.close()

```

Listing 1: Script to extract token and bigram information

The top 10 tokens in terms of frequency are shown in Table 1. In file results.csv you can find the complete results for the tokens' frequency ranking in this collection. Table 2 shows the same information for the case of the bigrams as can be also found in results\_bigrams.csv. To generate the requested graphs I used gnuplot on the results.csv, results\_bigrams.csv files. The log-log rank-frequency plots for the tokens, bigrams and their comparison are presented in Figures 1, 2 and 3, respectively.

Table 1: Top ten ranked tokens in Wiki small collection

| Rank | Frequency | Word      | Probability           | c                    |
|------|-----------|-----------|-----------------------|----------------------|
| 1    | 169695    | the       | 0.038954823010881046  | 0.038954823010881046 |
| 2    | 111772    | of        | 0.025658142417703502  | 0.051316284835407004 |
| 3    | 77940     | and       | 0.01789174050778201   | 0.05367522152334603  |
| 4    | 62273     | a         | 0.014295257334373996  | 0.057181029337495984 |
| 5    | 60351     | Wikipedia | 0.01385404710527524   | 0.06927023552637619  |
| 6    | 59125     | in        | 0.013572609154767917  | 0.08143565492860749  |
| 7    | 54327     | to        | 0.012471190487121803  | 0.08729833340985262  |
| 8    | 41122     | is        | 0.009439878793443827  | 0.07551903034755061  |
| 9    | 33865     | by        | 0.0077739773196822915 | 0.06996579587714062  |
| 10   | 31434     | The       | 0.007215922133969974  | 0.07215922133969974  |

Table 2: Top ten ranked bigrams in Wiki small collection

| Rank | Frequency | Word                     | Probability           | c                    |
|------|-----------|--------------------------|-----------------------|----------------------|
| 1    | 37574     | ('of', 'the')            | 0.008625407465221982  | 0.008625407465221982 |
| 2    | 15151     | ('in', 'the')            | 0.003478031311693678  | 0.010434093935081035 |
| 3    | 14010     | ('is', 'a')              | 0.0032161057802672054 | 0.012864423121068821 |
| 4    | 12147     | ('the', 'free')          | 0.0027884394655892752 | 0.013942197327946377 |
| 5    | 12088     | ('free', 'encyclopedia') | 0.0027748955511684497 | 0.016649373307010697 |
| 6    | 12088     | ('Wikipedia', 'the')     | 0.0027748955511684497 | 0.019424268858179147 |
| 7    | 12086     | ('About', 'Wikipedia')   | 0.0027744364354253706 | 0.027744364354253707 |
| 8    | 10932     | ('Wikipedia', 'user')    | 0.0025095266516688857 | 0.03011431982002663  |
| 9    | 10932     | ('by', 'Wikipedia')      | 0.0025095266516688857 | 0.032623846471695514 |
| 10   | 9622      | ('user', 's')            | 0.002208805839952252  | 0.030923281759331532 |




Figure 1: The rank-frequency curve for the words in the Wiki small collection




Figure 2: The rank-frequency curve for the bigrams in the Wiki small collection




Figure 3: The rank-frequency curve for the bigrams and words in the Wiki small collection

## Question 4.2

Plot vocabulary growth for the Wikipedia collection and estimate the parameters for Heaps law. Should the order in which the documents are processed make any difference?

### Answer

For this assignment we should count the number of words per document and the number of unique words. This way we have a analogy between the total number of words and the number of unique words in different points steps of parsing the collection of documents. I used the Wiki small collection again. The script used is shown in listing 2. It can take as option argument the -r option that will make it parse the documents in reverse order for the purposes of this assignment. The parameters that need to be defined are k and  $\beta$  from Heap's law:

$$v = kn^\beta$$

In order to estimate those values, I used non-linear least square fitting from python. The script used for this purpose is presented in listing 3. The curves produced by using the values found from this fitting and the curves generated by the experiments are shown in figure 4 and figure 5. Using this methodology I get the results shown in table 3, which point out that the order in which the files are visited can impact the vocabulary growth curve.

```
import os
import csv
from bs4 import BeautifulSoup
import nltk
import argparse
6
8 parser = argparse.ArgumentParser('vocgrowth')
parser.add_argument('--reverse', '-r', help='Visit files in reverse order', default=False, type=bool)
10 args = parser.parse_args();
reverse = args.reverse
12 # Set the path to search for html files
root_dir = 'en'
14
# Set the tokenizer expression
16 tokenizer = nltk.RegexpTokenizer(r'\w+')
18 # Find all files we need to parse
html_files = []
20 for root, dirs, files in os.walk(root_dir):
    for file in files:
        if file.endswith('.html'):
            html_files.append(os.path.join(root, file))
24
if reverse:
    html_files.reverse()
# Get the tokens for all files
28 tokens_found = []
voc_growth = []
30 i = 0;
for file in html_files:
    i+=1
    with open(file) as curr_file:
        print file
        soup = BeautifulSoup(curr_file.read(), 'html.parser')
36        tokens= tokenizer.tokenize(soup.get_text())
        for token in tokens:
            token = token.encode('utf-8')
            tokens_found.append(token)
38
```

```

40     voc_growth.append([i ,len(tokens_found),len(set(tokens_found))])
41
42 # Write token results to a file
43 if reverse:
44     f = open('voc_growth_reverse.csv','w+')
45 else:
46     f = open('voc_growth.csv','w+')
47 writer = csv.writer(f)
48 writer.writerow(['Files','Words','Vocabulary_size'])
49 for v in voc_growth:
50     writer.writerow( (v[0],v[1],v[2]))
51
52 f.close()

```

Listing 2: Script to extract the vocabulary growth information

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy.optimize import leastsq
4 import csv
5
6 f = open('voc_growth_reverse.csv','r') # Or voc_growth.csv :in order parsed files
7 i=0
8 xdata = []
9 ydata = []
10 for line in f:
11     if(i==0):
12         i=1
13         continue
14     parts = line.split(',')
15     xdata.append(int(parts[1]))
16     ydata.append(int(parts[2].strip()))
17 f.close()
18 # print len(xdata)
19 # print len(ydata)
20 xx= np.array(xdata)
21 yy= np.array(ydata)
22 p0 = np.array([1,1])
23
24 def my_fun(par,x,y):
25     return y-par[0]*np.power(x,par[1])
26
27 def fun_eval(par,x):
28     return par[0]*np.power(x,par[1])
29
30 # print 1*p0[0]
31 result = leastsq(my_fun,p0[:,],args=(xx,yy))
32 f = open('fit.csv','w+')
33 writer = csv.writer(f)
34 for x in xx:
35     writer.writerow((x,fun_eval(result[0],x)))
36 k,b = result[0]
37 print "k = "+str(k),
38 print "b = "+str(b)
39 f.close()

```

Listing 3: Script that calculates the parameters of Heap's law using non-linear least square fitting

Table 3: Heap's law parameters

|         | In order | Reverse |
|---------|----------|---------|
| k       | 25.1230  | 6.807   |
| $\beta$ | 0.597    | 0.680   |




Figure 4: The vocabulary growth for the Wiki small collection

Figure 5: The vocabulary growth for the Wiki small collection parsed in reverse order

Table 4: Comparison between the two stemming algorithms for the five html files chosen

| No | File                       | Porter stems # | Krovetz stems # |
|----|----------------------------|----------------|-----------------|
| 1  | Agia_Dynati_f24a.html      | 230            | 227             |
| 2  | Agnes_of_Glasgow_3312.html | 346            | 342             |
| 3  | Agoranomos.html            | 454            | 450             |
| 4  | Astolfo.html               | 628            | 632             |
| 5  | Parand.html                | 738            | 746             |

## Question 4.6

Process five Wikipedia documents using the Porter stemmer and the Krovetz stemmer. Compare the number of stems produced and find 10 examples of differences in the stemming that could have an impact on ranking.

### Answer

The five documents I chose for this assignment are shown below in table 4 with the number of stems per algorithm. I use existing Python's libraries to retrieve the results of both stemming algorithms. Krovetz's algorithm is more sophisticated and this can be realized by the 10 differences presented in table 5. As one can notice, Porter stemming produces many stems that have no meaning in English, thus, impacting the overall ranking of the document. The script used to collect this information is given in listing 4. This script produces 2 files per original document, containing the Porter and Krovetz stemming respectively. The output can be found in my github repository in filenames ending in .krovetz or .porter.

```

1 import io
2 import sys
3 import os
4 from bs4 import BeautifulSoup
5 import nltk
6 from nltk.stem import *
7 import krovetzstemmer

9 tokenizer = nltk.RegexpTokenizer(r '\w+')
10 porterizer = PorterStemmer()
11 krovetzizer = krovetzstemmer.Stemmer()

13 print('----- Starting -----')

```

```

15 if len(sys.argv)< 6:
16     print('Please provide 5 files to process')
17     print('porter_krovetz <file1> <file2> <file3> <file4> <file5>')
18     exit()
19
20 files = []
21 for arg in sys.argv[1:]:
22     files.append(arg)
23
24 porter_stems = []
25 krovetz_stems = []
26 header = ('No', 'File', 'Porter stems #', 'Krovetz stems #')
27 results = []
28 i=0
29 print header
30 for file in files:
31     i+=1
32     with open(file) as curr_file:
33         # print file
34         soup = BeautifulSoup(curr_file.read(), 'html.parser')
35         tokens= tokenizer.tokenize(soup.get_text())
36         porter = []
37
38         f = open(os.path.split(file)[1], 'w')
39         p = open(os.path.split(file)[1]+'.porter', 'w')
40         k = open(os.path.split(file)[1]+'.krovetz', 'w')
41         words = 0;
42         for token in tokens:
43             if(token.isalnum()):
44                 words+=1
45                 f.write(token.lower().encode('utf-8')+ ' ')
46                 porter_stems.append(porterizer.stem(token))
47                 krovetz_stems.append(krovetzizer.stem(token))
48                 p.write(porterizer.stem(token).encode('utf-8')+ ' ')
49                 k.write(krovetzizer.stem(token).encode('utf-8')+ ' ')
50                 if(words%15 == 0):
51                     f.write('\n')
52                     p.write('\n')
53                     k.write('\n')
54         results.append((i,os.path.split(file)[1],len(set(porter_stems)),len(set(krovetz_stems))))
55
56 print results[i-1]

```

Listing 4: Script to extract stems from the 5 documents using the two given algorithms

Table 5: Ten differences in the stems produced by the 2 algorithms

| Original | Porter  | Krovetz  |
|----------|---------|----------|
| produces | produc  | produce  |
| believe  | believ  | believe  |
| greece   | grec    | greece   |
| because  | becaus  | because  |
| his      | hi      | his      |
| zeus     | zeu     | zeus     |
| during   | dure    | during   |
| infamous | infam   | infamous |
| massacre | massacr | massacre |
| division | divis   | division |

Table 6: Top ten documents with the most inlinks with their anchors

| No | Inlinks # | Document   | Anchor texts   |
|----|-----------|--|--|
| 1  | 123       | en/articles/b/r/a/Brazil.html                              | Brazil,BRA,Brazilian   |
| 2  | 44        | en/articles/a/u/g/August_26.html                           | 08-26,26,August 26,26 August                                 |
| 3  | 17        | en/articles/m/a/n/Manga.html                               | manga,Manga  |
| 4  | 14        | en/articles/m/a/g/Magazine.html                            | magazine,magazines,Magazine                                  |
| 5  | 12        | en/articles/m/o/l/Mollusca.html                            | Mollusca   |
| 6  | 11        | en/articles/v/i/c/Victoria_of_the_United_Kingdom_5e8e.html | Victoria,Queen,Victoria of the United Kingdom,Queen Victoria |
| 7  | 8         | en/articles/s/c/r/Screenwriter.html                        | Writer(s),screenwriter,Screenwriter                          |
| 8  | 7         | en/articles/t/u/s/Tuscany.html                             | Tuscany  |
| 9  | 6         | en/articles/k/i/d/Kidney.html                              | kidneys,Renal,kidney   |
| 10 | 6         | en/articles/t/o/t/Tottenham_Hotspur_F.C._6bd2.html         | Tottenham Hotspur,Tottenham                                  |

## Question 4.8

Find the 10 Wikipedia documents with the most inlinks. Show the collection of anchor text for those pages.

## Answer

The methodology for this assignment is to parse all files in the Wiki small collection, maintaining a dictionary of documents to number of inlinks and another of inlinks to anchor texts. Once all files have been parsed, the first dictionary is sorted by the number of inlinks and then the second dictionary is printed based on the order that the first dictionary has given to its keys. Only the links that point to file in the collection are considered for this assignment. The results of the algorithm can be presented in table 6. The Python script for this assignment is presented in listing 5

```
1 import os
2 import csv
3 import operator
4 from bs4 import BeautifulSoup
5 import nltk
6
7 root_dir = 'en'
8
9 # Set the tokenizer expression
10 tokenizer = nltk.RegexpTokenizer(r'\w+')
11
12 # Find all files we need to parse
13 html_files = []
14 for root, dirs, files in os.walk(root_dir):
15     for file in files:
16         if file.endswith('.html'):
17             html_files.append(os.path.join(root, file))
18
19 file_str = []
20 for file in html_files:
21     file_str.append(str(file))
22
23 inlinks = {}
24 anchor_texts = {}
25
26 for file in html_files:
27     print file
28     soup = BeautifulSoup(open(file), 'html.parser')
29     links = soup.find_all('a')
30     for link in links:
31         new_link = link.get('href')
32         if not new_link:
33             continue
34         if new_link.startswith('http'):
35             continue
36         new_link = new_link.replace('../', '')
37         new_link = new_link.encode('utf-8')
38         new_link = 'en/' + new_link
```

```

37     if new_link in file_str and file != new_link:
38         if new_link not in inlinks:
39             inlinks[new_link] = 0
40             anchor_texts[new_link] = set()
41             inlinks[new_link] += 1
42             anchor_texts[new_link].add(link.text)
43 i=0
44 f = open("Inlinks_count", "w")
45 writer = csv.writer(f)
46 writer.writerow(( 'No' , 'Document' , 'Inlinks #' , 'Anchor texts'))
47 for k,v in sorted(inlinks.items(), key=operator.itemgetter(1), reverse=True):
48     i+=1
49     print i, str(v), k,
50     anchors = ','
51     for it in anchor_texts[k]:
52         anchors+=it+','
53     writer.writerow((i, str(v), k, anchors))
54     if(i==10):
55         break;
56 f.close()

```

Listing 5: Script to extract the top 10 documents with most inlinks and their anchor texts

## Question 5.8

Write a program that can build a simple inverted index of a set of text documents. Each inverted list will contain the file names of the documents that contain that word. Suppose the file A contains the text the quick brown fox, and file B contains the slow blue fox. The output of your program would be:

```
% ./your-program A B
blue B
brown A
fox A B
quick A
slow B
the A B
```

## Answer

In order to create the inverted index, I parse all the documents given, maintaining a dictionary of tokens to sets of files. For each token I keep all the documents that contain this term. The script used takes as command argument a list of documents to parse, or a filepath to search for all .html files below it and creates a file containing the inverted index. A part of the file produced is shown in table 7. I used the files Enayetpur.html,Gabrias.html for this assignment. The script that generates the inverted index is presented in listing 6

```

import os
import csv
from bs4 import BeautifulSoup
import nltk
import argparse
tokenizer = nltk.RegexpTokenizer(r'\w+')
parser = argparse.ArgumentParser('invindex')
group = parser.add_mutually_exclusive_group(required= True)
group.add_argument('-l', '--list', nargs='+', help='List of documents to parse',
required=False)

```

```

12 group.add_argument(' -p ', '--path', help='Path to directory to parse all of its files',
13                         required=False)
14 args = parser.parse_args()
15
16 html_files = []
17 if(args.path):
18     root_dir = args.path
19     for root,dirs,files in os.walk(root_dir):
20         for file in files:
21             if file.endswith('.html'):
22                 html_files.append(os.path.join(root, file))
23 else:
24     html_files = [x for x in args.list if x.endswith('.html')]
25
26 inv_index = {}
27
28 for file in html_files:
29     with open(file) as f:
30         soup = BeautifulSoup(f.read(), 'html.parser')
31         tokens= tokenizer.tokenize(soup.get_text())
32         for token in tokens:
33             token = token.encode("utf-8")
34             if token not in inv_index:
35                 inv_index[token] = set()
36             inv_index[token].add(os.path.split(file)[1])
37             # break;
38
39 outfile = open("inv_file.csv","w")
40 for token in inv_index:
41     outfile.write(token+",")
42     i=0
43     outfile.write("\n")
44     for file in inv_index[token]:
45         i+=1
46         outfile.write(file)
47         if(i != len(inv_index[token])):
48             outfile.write(",")
49     outfile.write("\n")

```

Listing 6: Script to create an inverted index

Table 7: The 30 first rows of the inverted index produced

| Token         | Documents                   |
|---------------|-----------------------------|
| Based         | Enayetpur.html,Gabrias.html |
| help          | Enayetpur.html,Gabrias.html |
| just          | Enayetpur.html              |
| text          | Enayetpur.html,Gabrias.html |
| Coordinates   | Gabrias.html                |
| produced      | Enayetpur.html              |
| sources       | Enayetpur.html              |
| geography     | Gabrias.html                |
| adding        | Enayetpur.html              |
| Recently      | Enayetpur.html              |
| also          | Enayetpur.html              |
| Navigation    | Enayetpur.html,Gabrias.html |
| 7EMonobook    | Enayetpur.html,Gabrias.html |
| YUNUS         | Enayetpur.html              |
| Here          | Enayetpur.html              |
| thana         | Enayetpur.html              |
| to            | Enayetpur.html,Gabrias.html |
| window        | Enayetpur.html,Gabrias.html |
| Alaibot       | Gabrias.html                |
| Nederlands    | Gabrias.html                |
| rich          | Enayetpur.html              |
| under         | Enayetpur.html,Gabrias.html |
| SieBot        | Gabrias.html                |
| main          | Enayetpur.html,Gabrias.html |
| geographical  | Gabrias.html                |
| lacking       | Enayetpur.html              |
| Khukni        | Enayetpur.html              |
| LordAnubisBOT | Gabrias.html                |
| material      | Enayetpur.html              |
| deductible    | Enayetpur.html,Gabrias.html |